

An Integrative Account of Constraints on Cross-Situational Learning

Daniel Yurovsky* and Michael C. Frank*

*Department of Psychology, Stanford University

Submitted to Proceedings of the National Academy of Sciences of the United States of America

Word-object co-occurrence statistics are a powerful information source for vocabulary learning, but there is considerable debate about learners actually use them. While some theories hold that we accumulate graded, statistical evidence about multiple referents for each word, others suggest that we track only a single candidate referent. In two large-scale experiments, we show that neither account is sufficient: Cross-situational learning involves elements of both. Further, the empirical data are captured by a computational model that formalizes how memory and attention interact with co-occurrence tracking. Together, the data and model unify opposing positions in a complex debate and underscore the value of understanding the interaction between computational and algorithmic levels of explanation.

statistical learning | word learning | language acquisition

Natural languages are richly structured. From sounds to phonemes to words to referents in the world, statistical regularities characterize the units and their connections at every level. Adults, children, and even infants are sensitive to these statistics, leading to a view of language acquisition as a parallel, possibly implicit, process of statistical extraction (1, 2). Recent experiments across a number of domains, however, show that human statistical learning may be significantly more limited than previously believed (3, 4, 5).

We focus here on the use of statistical regularities to learning the meanings of concrete nouns (known as cross-situational word learning; 6, 7, 8). Because word meanings are reflected in the statistics of use across contexts, learners could discover the meaning of the word “ball” (for instance) by noticing that while it is heard across many ambiguous contexts, it often accompanies play with small, round toys. A growing body of experiments shows that adults, children, and infants are sensitive to such co-occurrence information, and can use it to map words to their referents (8, 9, 10, 11).

Information about a word’s meaning can thus be extracted from the environmental statistics of its use (12). But this analysis is posed at what Marr (13) called the “computational theory” level: dealing only with the nature of the information available to the learner. At the “algorithmic” level—the level of psychological instantiation in the mind of the learner—this idealized statistical computation could be realized in many ways, and the computation human learners actually perform is a topic of significant debate.

Do human learners really maintain a representation of word-object co-occurrences? Some evidence suggests that humans are indeed gradual, parallel accumulators of statistical regularities about the entire system of word-object co-occurrences, simultaneously acquiring information about multiple candidate referents for the same word (14, 15, 16). Other evidence suggests that statistical learning is a focused, discrete process in which learners maintain a single hypothesis about the referent of any given word. This referent is either verified by future consistent co-occurrences or instead rejected, “re-setting” the learning process (17, 5). While both of these algorithmic-level solutions will, in the limit, produce successful word-referent mapping, they will do so at very different rates. In particular, if learners track a only a single referent for each word, it may be necessary to posit additional biases and constraints on learners in order for human-scale lexicons to be learned in human-scale time from the input available to children (18, 19).

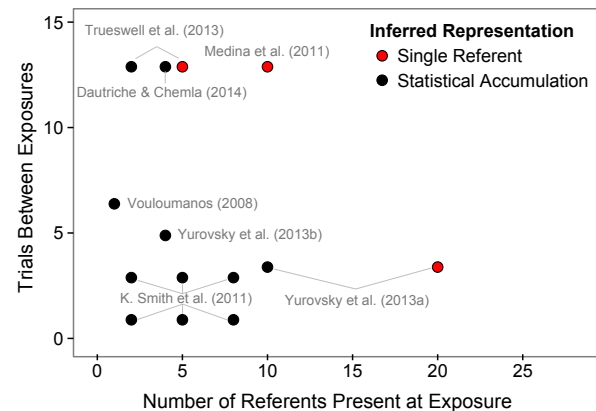


Figure 1. Results of previous experiments investigating representations for cross-situational learning. These experiments vary along a number of dimensions, but two appear to predict whether multiple-referent tracking is observed: the number of referents present on each trial, and the interval between trials for the referent.

To distinguish between these two accounts, previous experiments exposed learners to words and objects in which co-occurrence frequencies indicated several high-probability referents for the same word. At the group level, participants in these experiments showed gradual learning of multiple referents for the same word (e.g., 14, 20); but gradual, parallel learning curves can be observed at the group level even if individuals are discrete, single-referent learners (21, 17). Experiments measuring the same learner at multiple points—a stronger test—have produced mixed results. In some cases, learners showed clear evidence of tracking multiple referents for each word, suggesting a distributional approximation mechanism at the algorithmic level (22, 23, 24). In other experiments, however, learners appear to track only a single candidate referent, and to restart from scratch if their best guess is wrong (17, 5).

These mixed results expose a fundamental gap in our understanding of the mechanisms humans use to encode and track environmental statistics critical for learning language. Evidence for each account is separately compelling, but neither account can explain the evidence used to support the other. Because previous experiments differ along a number of dimensions—e.g., methodology, stimuli, timing, and precision of measurement—it has been difficult to integrate them to

Reserved for Publication Footnotes

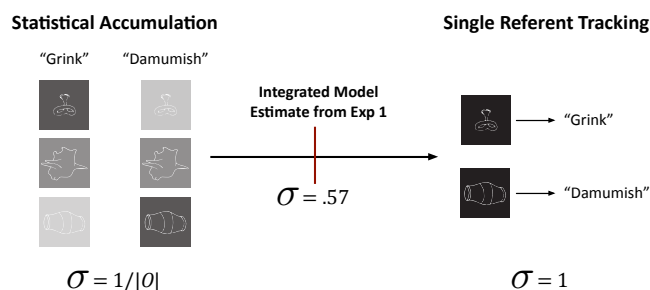


Figure 2. A representation of the continuum between the statistical accumulation and single-referent tracking models as learners' attention is varied from evenly distributed ($\sigma = \frac{1}{|O|}$) to focused on a single referent ($\sigma = 1$), as well as the best-fitting integrated model's position along this continuum.

understand why cross-situational learning sometimes appear distributional and sometimes appear discrete (for review, see 16).

We propose that differences in task difficulty may explain diverging results across experiments. Two salient dimensions vary across previous studies: ambiguity of individual learning instances, and the interval between successive exposures to the same label (Fig. 1). As attentional and memory demands increase, learners may shift from statistical accumulation to single-referent tracking (22, 5).

We present a strong test of this hypothesis, adapting a paradigm first introduced in (25) to study the information learners store in concept identification. We parametrically manipulated both the ambiguity of individual learning trials and the interval between them and measured multiple-referent tracking at the individual-participant level. Even at the maximum difficulty tested, learners tracked multiple referents for each word: Strong evidence against a qualitative shift from statistical accumulation to single-referent tracking. However, the data also show that learners encode the referents with differing strengths, remembering their hypothesized referent much better. Thus, each previous account appears to be partially correct.

To clarify how these two accounts are related, we implemented both the single-referent tracking and statistical accumulation as computational models. We also extended these accounts into an integrative model that subsumes both as special cases along a continuum. Only the integrative model accounted for our full dataset. This model then was able to make nearly perfect parameter-free predictions for a follow-up experiment that was designed to verify that learners encode mappings rather than individual words and objects. We conclude that cross-situational word learning is best characterized by an integrative account: Learners track both a single target referent and an approximation to the co-occurrence statistics; the strength of this approximation varies with the complexity of the learning environment (Fig. 2).

Experiment 1

We designed Experiment 1 to estimate learners' memory for both their single best hypothesis about the correct referent of a novel word and their additional statistical knowledge as demands on attention and memory varied. Participants saw a series of individually ambiguous word learning trials in which they heard one novel word, viewed multiple novel objects, and made guesses about which object went with each word. To succeed, participants needed to encode at least one of the objects that co-occurred with a word, remember it until their next encounter with that word, and check whether that same object was again present. If participants encoded exactly one object, they would succeed only when their initial hypothesis was correct. However, the more *additional* objects participants encoded on their first encounter with a word, the greater their likelihood of succeeding even if their initial hypothesis was incorrect.

Rather than allowing chance to determine whether participants held the correct hypothesis on their first exposure to a novel word,




Figure 3. A schematic of the experimental trials seen by participants in Experiments 1 and 2. On their first exposure to each novel word, participants were asked to guess its correct referent. In Experiment 1, the second trial for each word was either a *Same* trial—the set of referents contained the participant's previous hypothesis, or a *Switch* trial—the set of referents contained one the participant had previously *not* hypothesized. In Experiment 2, *Switch* trials were replaced with *New Label* trials that showed same set of referents but a played a novel word. The number of referents on the screen and the interval between successive exposures to the same word varied across conditions.

the set of novel objects presented on the second exposure to each word was constructed based on participants' choices. On *Same* trials, the participant's hypothesized referent was pitted against a set of novel competitors. In contrast, on *Switch* trials, one of the objects the participant had previously *not* hypothesized was pitted against a set of novel competitors (see Fig. 3). Logically, either a single-referent tracking or a statistical accumulation mechanism will succeed on *Same* trials. However, only statistical accumulation of information about non-target items can succeed on *Switch* trials.

Results. Do statistical learners encode multiple referents for each word, or do they instead encode only a single hypothesized referent? We compared the distribution of correct responses made by each participant to the distribution expected if participants were selecting randomly (defined by a Binomial distribution with four trials and a probability of success of $1/\#$ Referents). The top row of Fig. 4 shows participants' accuracies in identifying the referent of each word in all conditions for both kinds of trials (*Same* and *Switch*). At all Referent and Interval levels, both for *Same* and for *Switch* trials, participants' responses differed from those expected by chance (smallest $\chi^2(4) = 15.07$, all $ps < .01$). Thus, learners encode more than a single hypothesis in ambiguous word learning situations, even under high levels of memory and attentional load.

Next, to quantify the effect of each factor on word learning, we fit a mixed-effects logistic regression model to the data from the full dataset (26). This analysis showed significant main effects of Number of Referents, Interval, and Trial Type. In addition, the model showed a significant two-way interaction between Referents and Trial Type and a significant three-way interaction between all three factors (Table S1). Thus, while word learning was best at low levels of referential ambiguity and at low memory demands, the decreases in word learning observed on *Same* and *Switch* trials were due to different factors. For *Same* trials, the number of Referents played a relatively small role in the difficulty of learning, while the Interval between learning and test played a large role. However, for *Switch* trials, there was relatively little decline in word mapping as Interval increased but a large decline due to number of Referents.

These data suggest that neither the single-referent tracking nor the statistical accumulation account of cross-situational word learning is correct. Although learners did encode multiple referents, they did not encode them all with equal strength. Memory for the hypothesized referent was stronger than for non-hypothesized referents at all referent-set sizes and at all intervals. Further, the difference between




Figure 4. Proportion of repeated referents selected by participants at each combination of number of Referents and Interval on Same and Switch trials in Experiment 1, and Same and New Label trials in Experiment 2. Each datapoint represents 75 participants in Experiment 1 and 50 participants in Experiment 2. Error bars indicate 95% confidence intervals computed by non-parametric bootstrap. Learning in all conditions of Experiment 1 differed from chance and declined mostly due to Interval for Same trials but mostly due to Referents for Switch trials. Experiment 2 Same trials replicated performance in Experiment 1 Same trials, but New Label trials were different from Switch trials in all Referent and Interval conditions.

them grew with number of referents. Thus, it appears that a new account is necessary that integrates elements of both single-referent tracking and accumulative statistical tracking.

Before presenting a formal integrative account in the Model section below, we first rule out one other possibility. Because the set of foils for each target referent was distinct, participants could have succeeded on Switch trials by selecting the most familiar object regardless of which word they were hearing. If so, these data would be consistent with a slightly amended single-referent tracking account in which learners also have some residual memory for previously-seen objects but have not learned them as word-object mappings. Experiment 2 presents a new learning condition to test this possibility.

Experiment 2

Participants' above-chance accuracy on Switch trials in Experiment 1 provides evidence of their memory for multiple objects, but not necessarily for the formation of referential mappings between the objects and the novel words. To rule out this second possibility, Experiment 2 replaced Switch Trials with New Label trials in which participants saw an object they had previously *not* selected among a set of novel competitors but heard a *New Label* (Fig. 3). If success on Switch trials was due purely to referent familiarity, New Label trials should produce similar responses. In contrast, if success on Switch trials was due to a learned mapping between words and referents, New Label trials should show a different pattern of performance.

Results. Participants showed robust evidence of learning mappings (rather than simply tracking familiar objects). Whereas participants on Same trials were more likely than predicted by chance to select a referent they had previously seen but not guessed, participants in New Label trials were, in many cases, *less* likely than predicted by chance to select these same referents. Further, in all Referent and Interval conditions, performance on New Label trials differed significantly from performance on comparable Switch trials. That is, these

participants recognized these referents from their first exposure, and further recognized that they did not co-occur on their previous exposure with the label they heard at test (bottom row of Fig. 4).

In addition, a mixed-effects logistic regression largely reproduced the patterns observed in Experiment 1—word learning accuracies on Same trials declined predominantly due to Interval between learning and test, and very little due to the number of Referents. New Label trials were driven almost entirely by the number of Referents—as was the case with Switch trials in Experiment 1 (Table S2).

Taken together, these data are strong evidence that neither the single-referent tracking nor the statistical accumulation account of cross-situational word learning is correct. Instead, cross-situational word learning is best characterized by a combination of both of these mechanisms. In the next section, we formalize this idea.

Model

We begin by describing the computational-level learning problem posed by Experiment 1 using the model developed in (12). In this framework, the learner observes a set of situations S with the goal of determining the lexicon of word-object mappings L that produced them $P(L|S)$. We can use Bayes' rule to describe the inferential computation the learner must perform:

$$P(L|S) \propto P(S|L) P(L) \quad [1]$$

Each situation consists of two observed variables: objects (O) and words (W). In addition, situations implicitly contain an additional hidden variable: an intention (I) by the speaker to refer to one of the objects. Thus, speakers first choose an object from the set and then choose a referential label for it. The probability of a lexicon is given as the joint probability of observing all of the words, objects, and intentions given that lexicon, times the lexicon's prior probability:

$$P(L|S) \propto \prod_{s \in S} P(W_s, I_s, O_s, |L) P(L) \quad [2]$$

Because the referential intention mediates the relationship between words and objects (12), we can rewrite Equation 2 using the chain rule:

$$P(L|S) \propto \prod_{s \in S} P(W_s|I_s, L) P(I_s|O_s) P(L) \quad [3]$$

To make predictions from this model, we need to define the probabilities in Eq. 3. Following (12), we propose that the word (W) used to label the intended referent on each trial is chosen uniformly from the set of all words in the lexicon for that object (L_o). In addition, we propose a simple parsimony prior for the lexicon: A priori, the larger the set of words in the lexicon that refer to the same object O , the lower the probability of that lexicon: $P(L_o) \propto \frac{1}{|L_o|}$.

We can then take this computational-level description of the problem and add cognitive constraints to understand how the patterns observed in our data arise from the interaction of learning mechanisms, attention, and memory (see e.g., 27, 28). We start by describing how participants allocate their attention on each learning trial, a critical point of difference between the two different accounts of cross-situational learning.

In this framework, the most convenient place to integrate attention is in defining the learner's beliefs about $P(I|O)$, the probability of the speaker choosing to refer to each object in the set. One possibility is to let each object be equally likely to be the intended referent, implementing parallel Statistical Accumulation as in (12). Alternatively, the learner could place all of the probability mass on one hypothesized referent – implementing a Single Referent tracking strategy. A more flexible alternative is to assign some probability mass σ to the hypothesized referent, and divide the remainder evenly among the remaining objects: $\frac{1-\sigma}{|O|-1}$. This Integrated model subsumes the other two as special cases: At $\sigma = 1$, it is a Single Referent tracker, and at $\sigma = \frac{1}{|O|}$, it is a parallel Statistical Accumulator (Fig. 3).

There is some debate about the mechanisms that give rise to attentional limitations (e.g. 29). In our formulation, attention is treated as a continuous resource, but this choice is a matter of convenience rather than a theoretical commitment. For our purposes, the important question is to what extent attention is focused on the single target referent, and a continuous implementation allows parameter-estimation to answer this question.

Next, we model how learners' memories for observed situations decay over time. We follow previous memory researchers by formalizing memory for a lexical entry as a *power function* of the interval between successive exposures (30). As with attention-allocation, there a number of successful models of the underlying mechanisms that give rise to phenomena like the power-law observed in human memory (e.g., 31, 32). However, as with attention, the critical aspect for modeling this data is to be consistent with the broader dynamics of human memory, rather than with determining which model can best account for these dynamics. Accordingly, memory for lexical entry L_o decays according to a power function of time t in which γ scales the strength of initial encoding and λ defines the rate of decay.

$$M(L_o) = \gamma L_o t^{-\lambda} \quad [4]$$

Finally, we provide a choice rule describing how learners select among the objects on each test trial. We propose that learners choose the correct referent with probability proportional to their memory for its lexical entry, and otherwise choose randomly among the set of referents.¹ We use this rule because all of the foils on both Same and Switch trials were novel, and thus should have no trace in memory.

Table 1. Goodness of fit measures for models on Experiments 1 and 2

Model	Log Likelihood	BIC	E1 r^2	E1+2 r^2
Statistical Accumulation	-6565	-13115	0.33	0.66
Single Referent	-5950	-11884	0.83	0.77
Integrated	-5590	-11157	0.95	0.97

The Integrated model outperformed both accounts on all measures.

$$P(\text{Correct}) = M(L_o) + \frac{1 - M(L_o)}{|O|} \quad [5]$$

All three models—Statistical Accumulation, Single Referent, and Integrated—were fit to the data from Experiment 1 at the individual-participant level. While the Single Referent and Statistical Accumulation models capture some of the structure in the data, each leaves significant variance unexplained. The Single Referent Model cannot predict above-chance performance on Switch trials, and the Statistical Accumulation model cannot predict a difference between the Same and Switch trials. The Integrated model, however, predicts 95% of the variance in the data, and significantly outperforms the other models in BIC comparisons as well—a metric that trades off its superior performance against its one additional parameter (Table 1).

We can use the models presented above, with parameters estimated from Experiment 1, to make parameter-free predictions about the data observed in Experiment 2. As before, the Single Referent and Statistical Accumulation models predict some of the variance in the new data, but leave much unexplained. The Integrated model makes near-perfect predictions about the new data—including the New Label condition—explaining 98% of the combined variance in the data from Experiments 1 and 2 (Table 1). Fig. 5 presents model predictions for all experimental data. Taken together, Experiments 1 and 2 and the integrated model results thus provide strong evidence that learners track not only a single hypothesis for the most likely referent of a novel word, but also some approximation to distributional statistics—in particular, an approximation that becomes less precise as referential uncertainty increases.

General Discussion

For an ideal learner, word-object co-occurrence statistics contain a wealth of information about meaning. But how is this information used by human learners? One possibility is that learning is fundamentally statistical, and we gradually accumulate distributional information across situations. Another possibility, however, is that we track only a single, discrete hypothesis at any time. While each of these accounts has some support in prior work, neither is consistent with all of the extant data.

Our results here suggest a synthetic explanation: The degree to which learners represent statistical information depends on the complexity of the learning situation. When there are many possibilities, learners represent little about any other than the one that is currently favored; when there are few, learners represent more. This account does not depend on positing multiple, discrete learning systems. Instead, the tradeoff between the most likely hypothesis and the alternatives emerges from graded constraints on memory and attention. Consistent with this account, when we manipulated the cognitive demands of a cross-situational word learning paradigm, we found a gradual shift in the fidelity with which alternatives were represented.

The graded shift in representation we found was well-described by an ideal learning model, but only when this model was modified to take into account psychological constraints on attention and memory (34, 11, 16). This framework allowed us to estimate the effects of

¹This formulation is equivalent to using Luce's Choice Axiom (33) in which the target has strength $M(L_o) + \frac{1-M(L_o)}{|O|}$ and each foil has strength $\frac{1-M(L_o)}{|O|}$.




Figure 5. Predictions of the Integrated model for all conditions in Experiment 1 and 2. This model was able to account for 98% of the variance in the data, significantly outperforming both Single Hypothesis and Parallel Accumulation models.

these constraints on learning to find the model that best fit the data—one intermediate between the two extreme poles of parallel statistical accumulation and single-referent tracking. This unifying account provides a route by which both hypotheses and sensitivity to statistics can make complementary contributions to word learning (35, 36).

The shift from a computational to an algorithmic (psychological) description was critical in capturing the pattern of human performance in our task (13, 27, 4). For the current model, we chose one principled instantiation of cognitive limitations based on previous work, but there may be other consistent proposals. Indeed, recent work from (37) suggests that human performance observed in cross-situational learnings task can be consistent with a number of seemingly quite different models that can mimic each other (see also, 38). These authors note that modeling choices peripheral to the central learning mechanism—e.g., attentional allocation, memory, choice rule—can be varied to produce many different patterns of learning. In order to address this issue, we fit a large set of parametrically-varying data that imposes strong constraints on model parameters and modeling choices. In addition, we prevented overfitting by fixing model parameters using Experiment 1 and making parameter-independent predictions about learning that were supported in Experiment 2. This approach allowed us to gain insight about both the central learning mechanism and the constraining processes that together determine human performance.

Our work shows how ideal learning models can inform psychological accounts of statistical learning more broadly. Although we focused on noun learning, our results are relevant for many problems in language, including phonetic category learning, speech segmentation, and grammar learning. In each of these domains, researchers have debated the degree to which learners represent distributional information (39, 27, 40). We suggest a synthesis: Learning is fundamentally distributional, but the fidelity of learners' distributional estimates depends critically on their limited attention and memory.

Materials and Methods

Participants. Experiments 1 and 2 were posted to Amazon Mechanical Turk as a set of Human Intelligence Tasks (HITs) to be completed only by participants with US IP addresses that paid 30 cents each (for a detailed comparison of laboratory and Mechanical Turk studies see 41). Ninety HITs were posted for each of the 16 Referent \times Interval conditions in Experiment 1 for a total of 1440 paid HITs. Sixty HITs were posted for each of the sixteen Referent \times Interval conditions in Experiment 2 for a total of 960 paid HITs. If a participant completed the experiment more than once, he or she was paid each time but only data from the first HIT

completion was included in the final data set (excluded 180 HITs in Experiment 1, 100 HITs in Experiment 2). In addition, data was excluded from the final sample if participants did not give correct answers for familiar trials (64 HITs in Experiment 1 and 60 HITs in Experiment 2, see Design and Procedure). The final sample for Experiment 1 comprised 1,196 unique participants, approximately 75 participants per condition (range: 71–81). The final sample for Experiment 2 comprised 803 unique participants, approximately 50 participants per condition (range: 41–55).

Stimuli. Stimuli for the experiment consisted of black and white pictures of familiar and novel objects and audio recordings of familiar and novel words. Pictures of 32 familiar objects spanning a range of categories (e.g. squirrel, truck, tomato, sweater) were drawn from the set constructed by (42). Pictures of distinct but difficult to name novel objects were drawn from the set of 140 first used in (43). For ease of viewing on participants' monitors, pixel values for all pictures were inverted so that they appeared as white outlines on black backgrounds (see Fig. 3). Familiar words consisted of the labels for the familiar objects as produced by AT&T Natural VoicesTM (voice: Crystal). Novel words were 1–3 syllable pseudowords obeying the rules of English phonotactics produced using the same speech synthesizer.

Design and Procedure. Participants were exposed to a series of trials in which they heard a word, saw a number of objects, and were asked to indicate their guess as to which object was the referent of the word. After a written explanation of this procedure, participants were given four practice trials to introduce them to the task. On each of these trials, they heard a Familiar word and saw a line drawing of that object among a set of other familiar objects. On the first two trials, participants were asked to find the squirrel, and the correct answer was in the same position on each trial. On the next two trials, participants were asked to find the sweater, and the correct answer switched positions from the first to the second trial (in order to ensure that participants understood the on-screen position was not an informative cue to the correct target). These trials also served to screen for participants who did not have their audio enabled or who were not attending to the task.

After the Familiar trials, participants were informed that they would now hear novel words, and see novel objects, and that they should continue selecting the correct referent for each word. Participants heard each of the eight novel words twice, but the order in which words were presented and the number of objects seen on the screen varied across sixteen between-subjects conditions. Participants saw either 2, 3, 4, or 8 Referents on each trial, and the two trials for each word occurred either back-to-back, or were interleaved between trials for other words for an Interval of 1, 2, 3, or 8. In both Experiments, four of these follow-up trials were Same trials in which the referent that participants selected on the first

² It is possible that forcing participants to select an object on each trial could have changed their performance. However, control conditions from three previous experiments suggest that empirically this is not the case (17, 22, 5).

encounter with that word appeared amongst the set of objects. In Experiment 1, the other four were Switch trials in which one of the referents in the set was selected randomly from the objects a participant *did not* select on the previous exposure to that word. In Experiment 2, the other four trials were New Label trials in which set of candidate referents was the same as on Switch trials in Experiment 1, but the word was novel (Fig. 3). All other referents were completely novel on each trial. The number of referents on Familiar trials for each participant matched the number of referents they would see on Same and Switch trials.

Because participants performed this task over the internet, it was important to inform them that their click had been registered. Thus, a red dashed box appeared around the selected object for 1 second after their click was received. This box appeared whether or not the selected object was the “correct” referent.²

ACKNOWLEDGMENTS. We are grateful to Linda Smith, Erika Bergelson, Molly Lewis, Ann Nordemeyer, and all of the members of the Language and Cognition Lab for their feedback on this project. This work was supported by a NIH NRSA F32HD075577 to DY and a grant from the Merck Scholars Foundation to MCF.

References

- Saffran JR, Aslin RN, Newport EL (1996) Statistical learning by 8-month-old infants. *Science* 274:1926–1928.
- Gómez RL, Gerken L (2000) Infant artificial language learning and language acquisition. *Trends Cogn Sci* 4:178–186.
- Johnson EK, Tyler MD (2010) Testing the limits of statistical learning for word segmentation. *Developmental Sci* 13:339–345.
- Yurovsky D, Yu C, Smith LB (2012) Statistical speech segmentation and word learning in parallel: Scaffolding from child-directed speech. *Front Psychol* 3:374.
- Trueswell JC, Medina TN, Hafri A, Gleitman LR (2013) Propose but verify: Fast mapping meets cross-situational learning. *Cognit Psychol* 66:126–156.
- Pinker S (1989) *Learnability and cognition: The acquisition of argument structure*. (Cambridge, MA: MIT Press).
- Siskind JM (1996) A computational study of cross-situational techniques for learning word-to-meaning mappings. *Cognition* 61:39–91.
- Yu C, Smith LB (2007) Rapid word learning under uncertainty via cross-situational statistics. *Psychol Sci* 18:414–420.
- Smith LB, Yu C (2008) Infants rapidly learn word-referent mappings via cross-situational statistics. *Cognition* 106:1558–1568.
- Scott RM, Fischer C (2012) 2.5-year-olds use cross-situational consistency to learn verbs under referential uncertainty. *Cognition* 122:163–180.
- Vlach HA, Johnson SP (2013) Memory constraints on infants’ cross-situational statistical learning. *Cognition* 127:375–382.
- Frank MC, Goodman N, Tenenbaum J (2009) Using speakers’ referential intentions to model early cross-situational word learning. *Psychol Sci* 20:578–585.
- Marr D (1982) *Vision: A Computational Investigation into the Human Representation and Processing of Visual Information*. (New York, NY: W. H. Freeman).
- Vouloumanos A (2008) Fine-grained sensitivity to statistical information in adult word learning. *Cognition* 107:729–742.
- McMurray B, Horst JS, Samuelson LK (2012) Word learning emerges from the interaction of online referent selection and slow associative learning. *Psychol Rev* 119:831–877.
- Yurovsky D, Fricker DC, Yu C, Smith LB (2014) The role of partial knowledge in statistical word learning. *Psychon B Rev* 21:1–22.
- Medina TN, Snedeker J, Trueswell JC, Gleitman LR (2011) How words can and cannot be learned by observation. *Proc Natl Acad Sci USA* 108:9014–9019.
- Vogt P (2012) Exploring the robustness of cross-situational learning under Zipfian distributions. *Cognitive Sci* 36:726–739.
- Reisenauer R, Smith K, Blythe RA (2013) Stochastic dynamics of lexicon learning in an uncertain and nonuniform world. *Phys Rev Lett* 110:258701–258705.
- Yurovsky D, Yu C, Smith LB (2013) Competitive processes in cross-situational word learning. *Cognitive Sci* 37:891–921.
- Gallistel CR, Fairhurst S, Balsam P (2004) The learning curve: Implications of a quantitative analysis. *Proc Natl Acad Sci USA* 101:13124–31.
- Smith K, Smith ADM, Blythe RA (2011) Cross-situational learning: An experimental study of word-learning mechanisms. *Cognitive Sci* 35:480–498.
- Yurovsky D, Smith LB, Yu C (2013) Statistical word learning at scale: The baby’s view is better. *Developmental Sci* 16:959–966.
- Dautriche I, Chemla E (in press) Cross-situational word learning in the right situations. *J Exp Psychol Learn*.
- Bower G, Trabasso T (1963) Reversals prior to solution in concept identification. *J Exp Psychol* 66:409–418.
- Baayen RH, Davidson DJ, Bates DM (2008) Mixed-effects modeling with crossed random effects for subjects and items. *J Mem Lang* 59:390–412.
- Frank MC, Goldwater S, Griffiths TL, Tenenbaum JB (2010) Modeling human performance in statistical word segmentation. *Cognition* 117:107–125.
- Shi L, Griffiths TL, Feldman NH, Sanborn AN (2010) Exemplar models as a mechanism for performing Bayesian inference. *Psychol B Rev* 17:443–464.
- Wei Z, Wang XJ, Wang DH (2012) From distributed resources to limited slots in multiple-item working memory: A spiking network model with normalization. *J Neurosci* 32:11228–11240.
- Anderson JR, Schooler LJ (1991) Reflections of the environment in memory. *Psychol Sci* 2:396–408.
- Murdock BB (1982) A theory for the storage and retrieval of item and associative information. *Psychol Rev* 89:609–626.
- Shiffrin RM, Steyvers M (1997) A model for recognition memory: REM - retrieving effectively from memory. *Psychon B Rev* 4:145–166.
- Luce RD (1959) *Individual choice behavior: A theoretical analysis*. (New York, NY: Wiley).
- Kachergis G, Yu C, Shiffrin RM (2012) An associative model of adaptive inference for learning word-referent mappings. *Psychon B Rev* 19:317–324.
- Waxman SR, Gelman SA (2009) Early word-learning entails reference, not merely associations. *Trends Cogn Sci* 13:258–263.
- Kachergis G, Yu C, Shiffrin RM (2013) Actively learning object names across ambiguous situations. *Top Cogn Sci* 5:200–213.
- Yu C, Smith LB (2012) Modeling cross-situational word-referent learning: Prior questions. *Psychol Rev* 119:21–39.
- Townsend JT (1990) Serial vs. parallel processes: Sometimes they look like Tweedledum and Tweedledee but they can (and should) be distinguished. *Psychol Sci* 1:46–54.
- Endress AD, Scholl BJ, Mehler J (2005) The role of salience in the extraction of algebraic rules. *J Exp Psychol Gen* 134:406–419.
- McMurray B, Kovack-Lesh KA, Goodwin D, McEchron W (2013) Infant directed speech and the development of speech perception: Enhancing development or an unintended consequence? *Cognition* 129:362–378.
- Crump MJC, McDonnell JV, Gureckis TM (2013) Evaluating Amazon’s Mechanical Turk as a tool for experimental behavioral research. *PLOS ONE* 8:e57410.
- Snodgrass JG, Vanderwart M (1980) A standardized set of 260 pictures: Norms for name agreement, image agreement, familiarity, and visual complexity. *J Exp Psychol Learn* 6:174–215.
- Kanwisher N, Woods RP, Iacoboni M, Mazziotta JC (1997) A locus in human extrastriate cortex for visual shape analysis. *J Cogn Neurosci* 9:133–142.